

# Data Streams:

## *Bloom Filters*

### Mining Massive Datasets

Prof. Carlos Castillo — <https://chato.cl/teach>



Universitat  
Pompeu Fabra  
*Barcelona*

# Sources

- Mining of Massive Datasets (2014) by Leskovec et al. (chapter 4)
  - Slides [part 1](#), [part 2](#)
- Tutorial: [Mining Massive Data Streams](#) (2019) by Michael Hahsler

# Bloom filters

# Filtering a data stream

- Suppose we have a large set  $S$  of keys
- We want to filter a stream  $\langle \text{key}, \text{data} \rangle$  to let pass only the elements for which  $\text{key} \in S$
- Example: key is an e-mail address, we have a total of  $|S|=10^9$  allowed e-mail addresses

What's the Naïve solution?

# Filtering a data stream

- Suppose we have a large set  $S$  of keys
- We want to filter a stream  $\langle \text{key}, \text{data} \rangle$  to let pass only the elements for which  $\text{key} \in S$
- Example: key is an e-mail address, we have a total of  $|S|=10^9$  allowed e-mail addresses
- Naïve solution? Hash table won't work, too big!

# Bloom Filter (1-bit case)

- Given a set of keys  $S$
- Create a bit array  $B[]$  of  $n$  bits
  - Initialize to all 0s
- Pick a hash function  $h$  with range  $[0, n)$ 
  - For each member of  $s \in S$ 
    - Hash to one of  $n$  buckets
    - Set that bit to 1, i.e.,  $B[h(s)] \leftarrow 1$
- For each element  $a$  of the stream
  - Output  $a$  if and only if  $B[h(a)] == 1$

Bloom filter creation

Stream processing

# Bloom Filter is an approximate filter

- Can it output an element with a key not in  $S$ ?
- Can it not output an element with a key in  $S$ ?

# Bloom Filter is an approximate filter

- Can it output an element with a key not in  $S$ ?

Yes, due to hash collisions  $h(x)=h(y)$  when  $x \neq y$

- Can it not output an element with a key in  $S$ ?

No, because  $h(x)$  is always the same for  $x$

Bloom filters are *permissive* (not *strict*)



# Bloom filter

- A bloom filter is:
  - An array of  $n$  bits, initialized as 0
  - A collection of hash functions  $h_1, h_2, \dots, h_k$
  - A set  $S$  of  $m$  key values
- The purpose of the bloom filter is to allow all stream items whose key is in  $S$

# Bloom filter initialization

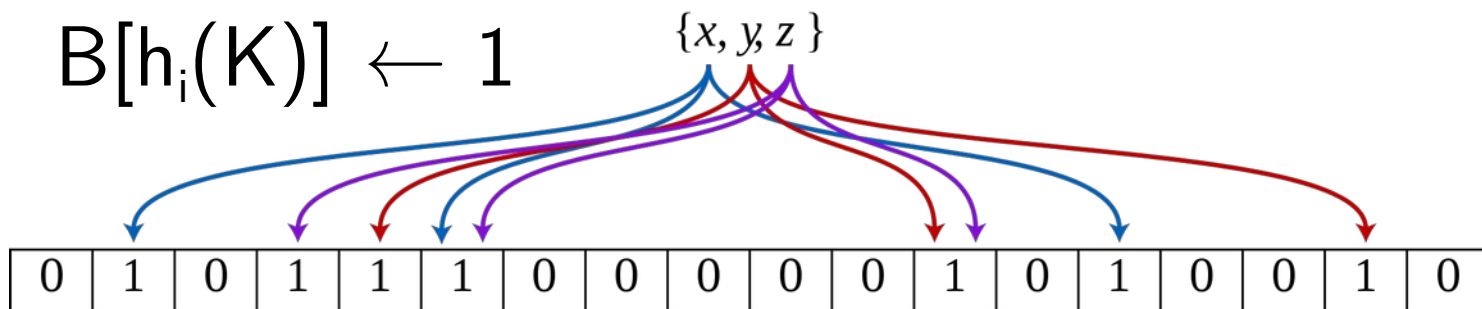
For all positions  $i$  in  $[0, n-1]$

$$B[i] \leftarrow 0$$

For all keys  $K$  in  $S$ :

For every hash function  $h_1, h_2, \dots, h_k$

$$B[h_i(K)] \leftarrow 1$$



# Bloom filter usage

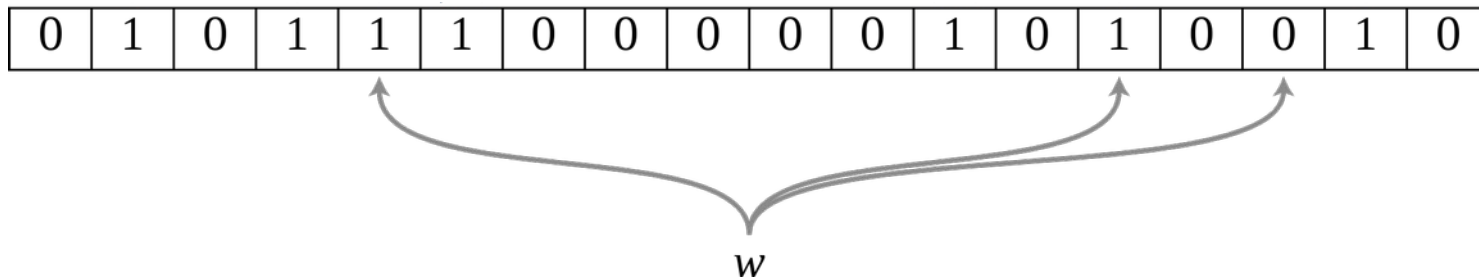
For each input element  $\langle \text{key}, \text{data} \rangle$

$\text{allow} \leftarrow \text{TRUE}$

For every hash function  $h_1, h_2, \dots, h_k$

$\text{allow} \leftarrow \text{allow} \wedge B[h_i(K)] == 1$

output element if  $\text{allow} == \text{TRUE}$



# Characteristics of Bloom Filters

- Are lax (not strict) and let some items pass
  - May require a second-level check to make filter strict, for instance store output on disk files and then check against hash tables (slower)
- Implementations can be very fast
  - E.g., use hardware words for the bit table

# Preliminaries for the analysis: targets and darts

- Suppose we throw  $y$  darts at  $x$  targets
  - All darts will hit one of the targets



$y=4$  darts

$x=4$  targets

# Preliminaries for the analysis: targets and darts (cont.)

- How many distinct targets can we expect to hit at least once?
  - Prob. that a given dart will hit a specific target is  $1/x$
  - Prob. that a given dart will **not** hit a specific target is  $1-1/x$
  - Prob. none of the  $y$  darts will hit a specific target is  $(1-1/x)^y = (1-1/x)^{x(y/x)}$
  - Using that  $(1-\epsilon)^{1/\epsilon} \simeq 1/e$  for small  $\epsilon$
  - If  $x$  is large,  $1/x$  is small, and prob. that none of the  $y$  darts will hit a specific target is  $(1/e)^{y/x}$



$y=4$  darts

$x=4$  targets

# Analysis of the 1-bit Bloom Filter

- Each element of the signature  $S$  is a dart  $|S|=y$
- Each bit in the array is a target  $n=x$
- Suppose  $y=|S|=10^9$  (1 G) and  $x=n=8 \times 10^9$  (8 G)
- Prob. that a given bit is **not** set to 1 (dart does not hit the target) is  $(1/e)^{y/x} = (1/e)^{1/8}$
- Prob. that a given bit is set to 1 is  $1 - (1/e)^{1/8} = 0.1175$
- Expected number of bits that is set to 1 = 11.75% x 8GB



About 12% of bits are set to one in this Bloom Filter

this is also the false-hit probability in this case

# General case

- $|S|=m$  keys, array has  $n$  bits
- $k$  hash functions
- Targets  $x=n$ , darts  $y=km$
- Probability that a bit remains 0 is  $(1/e)^{km/n} = e^{-km/n}$
- Example:

We can pick  $k=n/m$  to obtain collision probability  $1/e = 37\%$



# Analysis of a 2-bit Bloom Filter

- Suppose  $|S|=10^9$  (1 G) and  $n=8 \times 10^9$  (8 GB)
- Suppose we use two hash functions
- Prob. that a given bit is NOT set to 1 (dart does not hit the target) is  $(1/e)^{y/x} = (1/e)^{1/4}$
- Prob. a bit is set to 1 is  $1 - (1/e)^{1/4}$
- Prob. two bits are set to 1 is  $(1 - (1/e)^{1/4})^2 = 0.0493$
- We have a false hit probability of about 5% with two hash functions, while the probability was about 12% with only one

# How many hash functions to use?


Too few: test is too unspecific. Too many: table becomes too crowded.

- $m = 1$  billion,  $n = 8$  billion

$$k = 1: (1 - e^{-1/8}) = 0.1175$$

$$k = 2: (1 - e^{-1/4})^2 = 0.0493$$

- What happens as we keep increasing  $k$ ?
  - “Optimal” value of  $k$ :  $n/m \ln(2)$
  - In our case: Optimal  $k = 8 \ln(2) = 5.54 \approx 6$
  - Error at  $k = 6$ :  $(1 - e^{-1/6})^2 = 0.0235$



# Summary

# Things to remember

- How to initialize a Bloom filter
- How to use a Bloom filter
- Proofs for 1-bit, 2-bit case

# Exercises for TT22-T26

- Mining of Massive Datasets (2014) by Leskovec et al.
  - Exercises 4.2.5
  - Exercises 4.3.4
  - Exercises 4.4.5
  - Exercises 4.5.6